

S&DS 265 / 565
Introductory Machine Learning

Course Wrap Up

Thursday, December 9

Yale

Reminders

- Assignment 7 out (neural nets and RL)
- New due date: Tuesday December 14
- Assignment 6 will be graded by today
- Quiz 4: Open until 1pm today
- Final exam, Dec 21 at 7pm (cumulative, 3 hours, cheat sheet, practice exam end of this week)

For today

- Addendum on RL
- Follow up from last class
- Examples of recent ML research
- Course summary
- Final exam

When does learning take place?

Recall from Bellman equation that y_t is an expectation.

Learning takes place when expectations are violated. The receipt of the reward itself does not cause changes.

A Neural Substrate of Prediction and Reward

Wolfram Schultz, Peter Dayan, P. Read Montague*

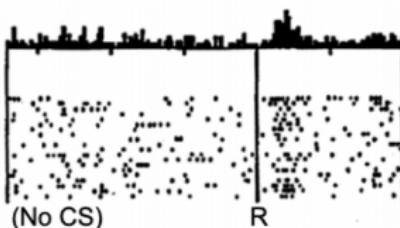
The capacity to predict future events permits a creature to detect, model, and manipulate the causal structure of its interactions with its environment. Behavioral experiments suggest that learning is driven by changes in the expectations about future salient events such as rewards and punishments. Physiological work has recently complemented these studies by identifying dopaminergic neurons in the primate whose fluctuating output apparently signals changes or errors in the predictions of future salient and rewarding events. Taken together, these findings can be understood through quantitative theories of adaptive optimizing control.

Science 1997

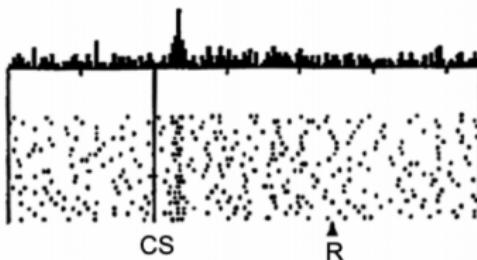
Neuroscience connection

Do dopamine neurons report an error
in the prediction of reward?

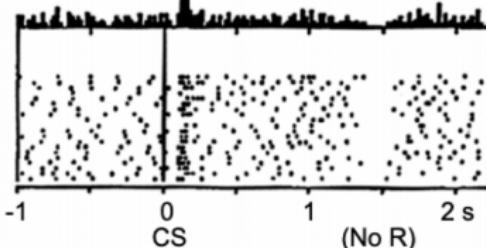
No prediction
Reward occurs



Reward predicted
Reward occurs



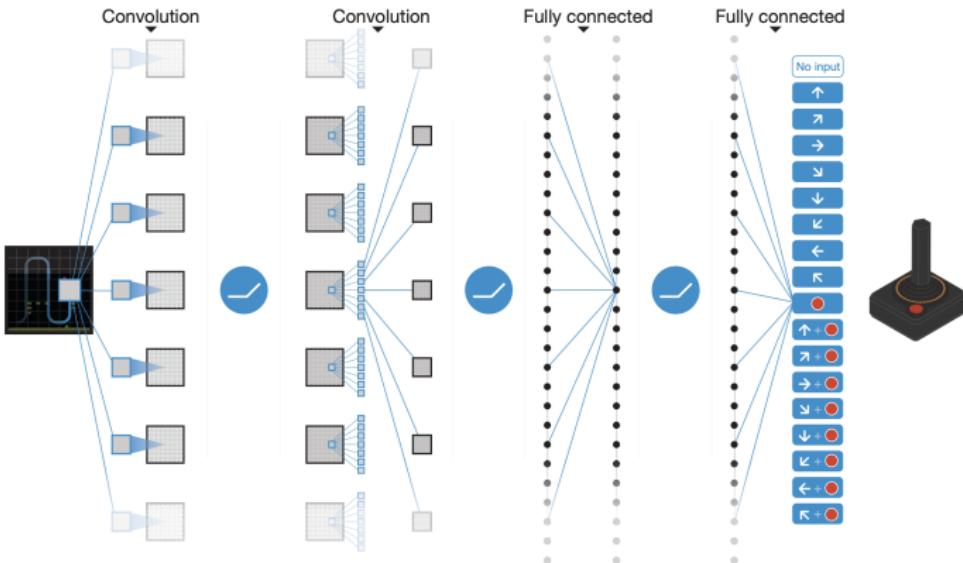
Reward predicted
No reward occurs



Space Invaders

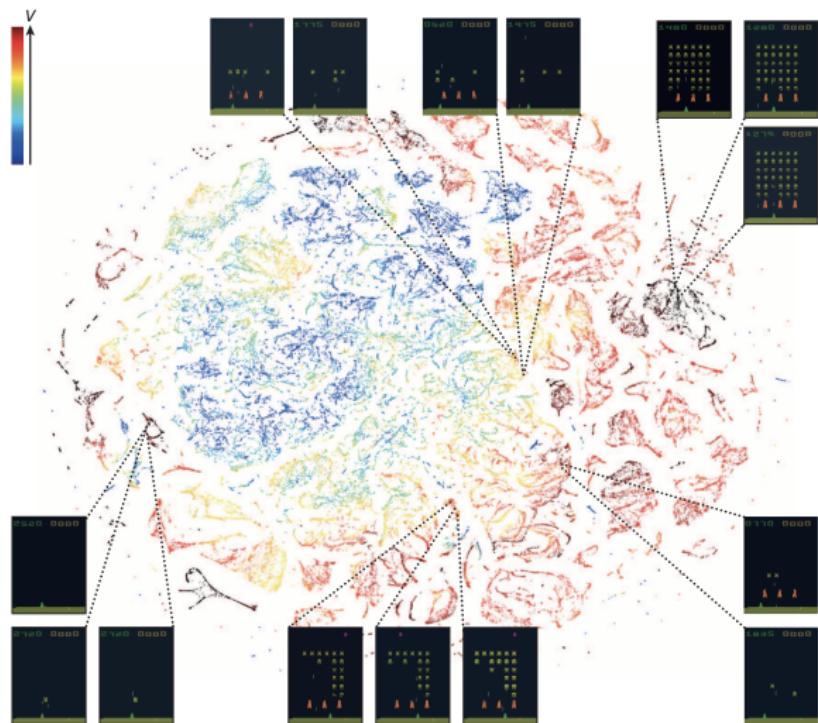


Second generation DQN



<https://storage.googleapis.com/deepmind-data/assets/papers/DeepMindNature14236Paper.pdf>

Second generation DQN: Interpretation



t-SNE representations of last layer for Space Invaders, color-coded for v_* .

Follow up from Tues



Is A.I. the Problem? Or Are We?

The Ezra Klein Show

Society & Culture

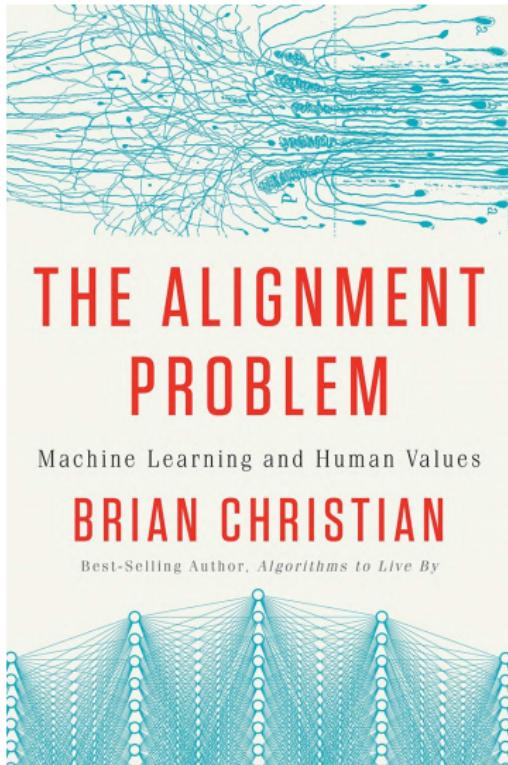
[Listen on Apple Podcasts ↗](#)



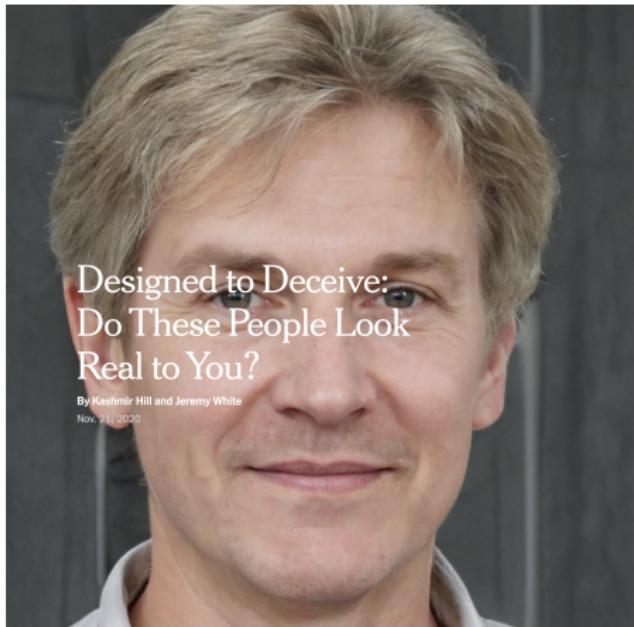
If you talk to many of the people working on the cutting edge of artificial intelligence research, you'll hear that we are on the cusp of a technology that will be far more transformative than simply computers and the internet, one that could bring about a new industrial revolution and usher in a utopia — or perhaps pose the greatest threat in our species's history.

Others, of course, will tell you those folks are nuts.

Follow up from Tues



Follow up from Tues



Designed to Deceive:
Do These People Look
Real to You?

By Kashmir Hill and Jeremy White

Nov. 21 | 2020

[https://www.nytimes.com/interactive/2020/11/21/science/
artificial-intelligence-fake-people-faces.html?](https://www.nytimes.com/interactive/2020/11/21/science/artificial-intelligence-fake-people-faces.html?)

Follow up from Tues



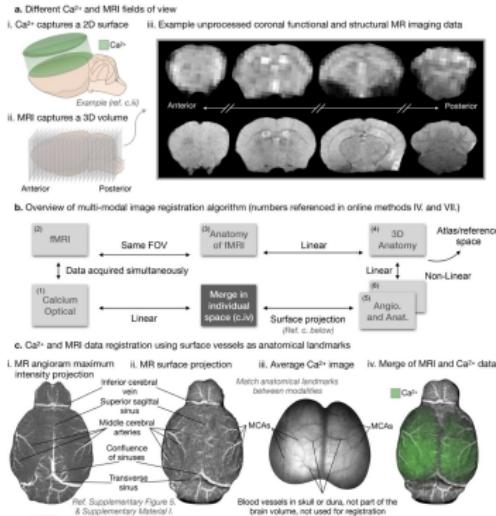
The image shows the header of the Yale Jackson Institute for Global Affairs website. It features a dark blue header bar with the institute's name in white. Below the header is a large banner with a globe in the center, set against a background of a circuit board pattern. A search bar and a menu icon are visible in the top left corner of the banner.

Jackson Institute establishes Schmidt Program on Artificial Intelligence, Emerging Technologies, and National Power

December 8, 2021 | by Yale Jackson

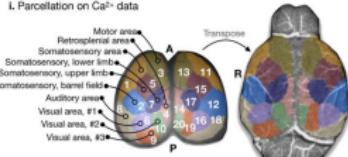
Some recent projects in Yale SML group

ML for brain imaging data

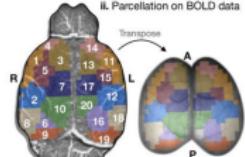


a. Representative Ca^{2+} and BOLD parcellation results transposed between Ca^{2+} and BOLD space

i. Parcellation on Ca^{2+} data

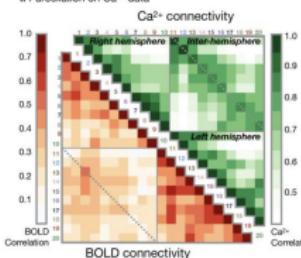


ii. Parcellation on BOLD data

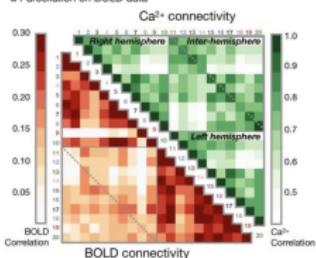


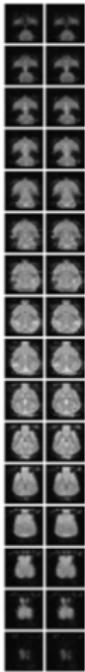
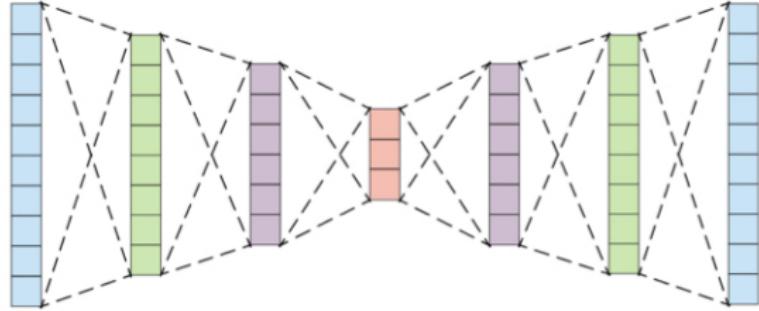
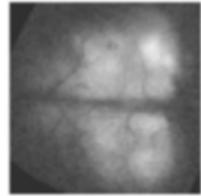
b. Average (N=6) Ca^{2+} and BOLD connectivity matrices

i. Parcellation on Ca^{2+} data



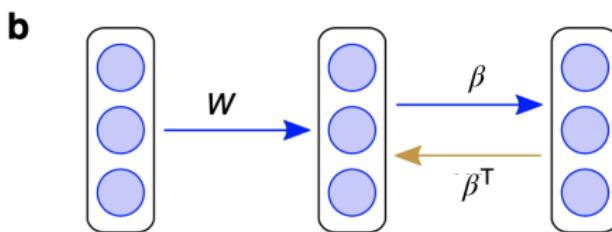
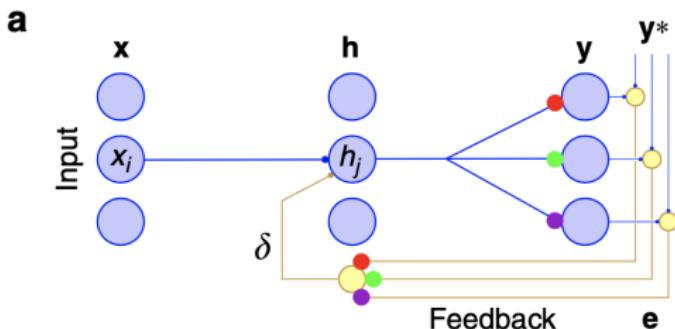
ii. Parcellation on BOLD data



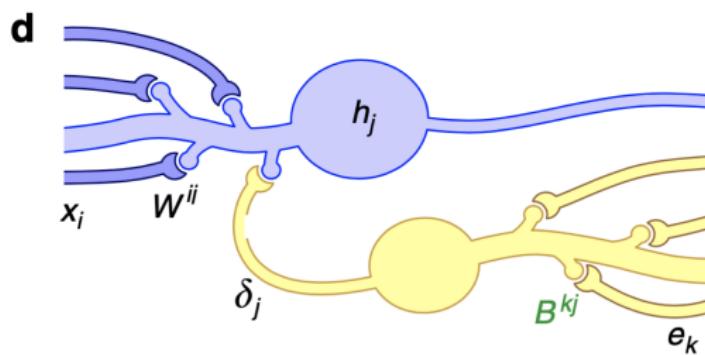
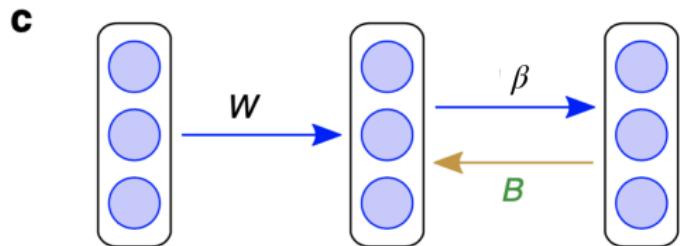


- U-net architecture with GANs
- 34 million parameters

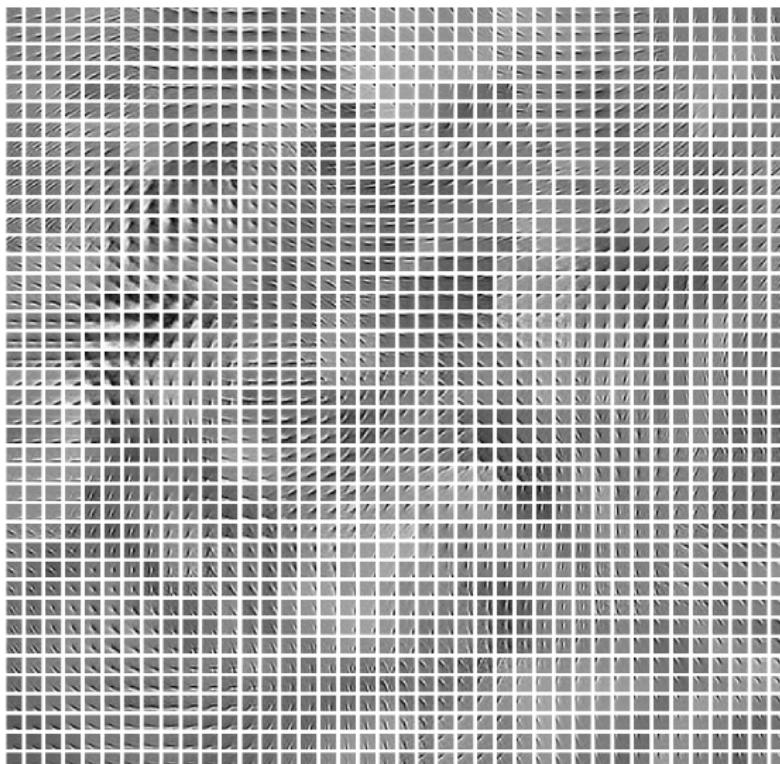
Proposal from DeepMind



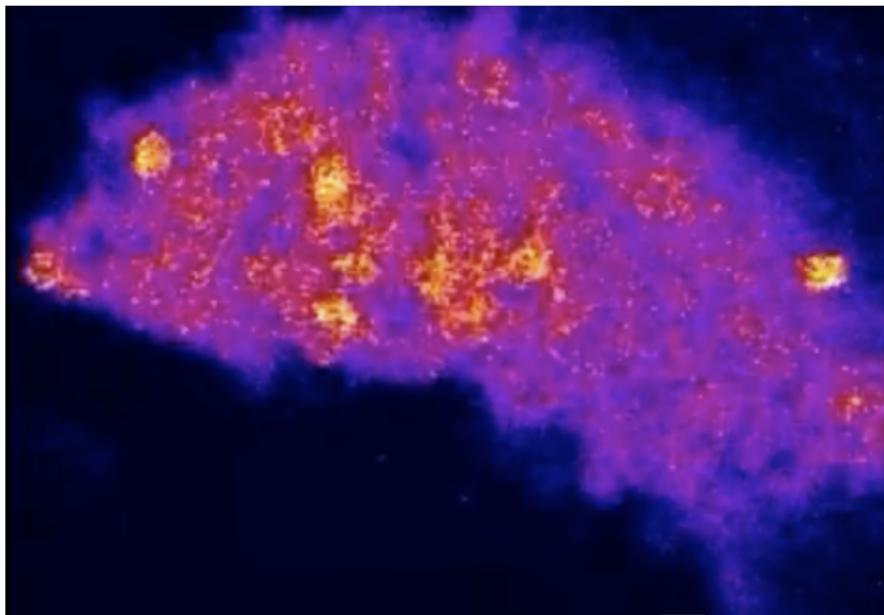
Proposal from DeepMind



Topographic maps of neural tunings

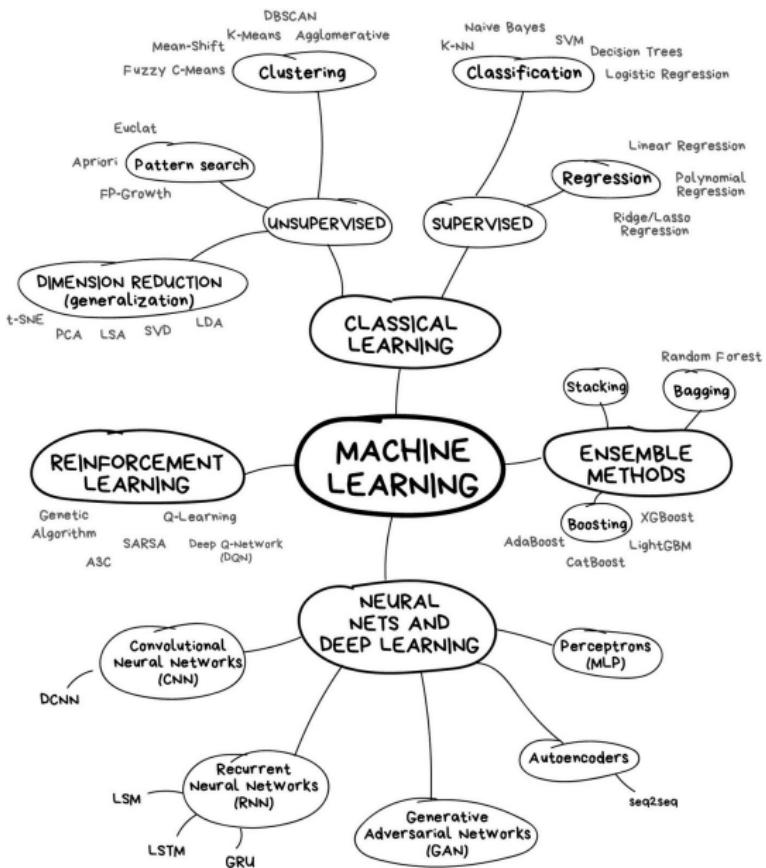


Sparse representations in memory systems



"Sparse, decorrelated odor coding in the mushroom body enhances learned odor discrimination," Lin et al., Nature Neuroscience (2014). Long line of work: Kanerva (1988), Marr (1969), Treves and Rolls (1991), Vinje and Gallant (2000), Hromadka et al. (2008), Crochet et al. (2011), Perez-Orive et al. (2002), Ito et al. (2008), Stettler and Axel (2009), Honegger et al. (2008).

We've covered a lot of ground!



3	Sept 14, 16	Linear regression and classification	<ul style="list-style-type: none">○ Covid trends (revised)○ Classification examples	Tue: <ul style="list-style-type: none">○ Assn1 out	Sept 16: Classification concepts Notes on regression Sept 16: Classification Notes on classification
4	Sept 21, 23	Stochastic gradient descent	<ul style="list-style-type: none">○ SGD examples	Tue: Quiz 1 Thu: Assn 1 in <ul style="list-style-type: none">○ Assn2 out	Sept 21: Classification (continued) Sept 23: Stochastic gradient descent
5	Sept 28, 30	Bias and variance, cross-validation	<ul style="list-style-type: none">○ Bias-variance tradeoff○ Covid trends (revised)○ California housing		Sept 28: Bias and variance Sept 30: Cross-validation
6	Oct 5, 7	Tree-based methods	<ul style="list-style-type: none">○ Trees and forests○ Visualizing trees	Tue: Assn 2 in <ul style="list-style-type: none">○ Assn3 out	Oct 5: Trees Oct 7: Forests
7	Oct 12, 14	PCA and dimension reduction	<ul style="list-style-type: none">○ PCA examples○ PCA revisited○ Used for regression	Tue: Quiz 2 Thu: Assn 3 in <ul style="list-style-type: none">○ Assn4 out	Oct 12: PCA Oct 14: PCA and review
8	Oct 19	Midterm exam (in class)			On Canvas: Practice midterm / Sample soln Midterm / Sample soln
9	Oct 26, 28	Language models, word embeddings	<ul style="list-style-type: none">○ Word embeddings		Oct 26: Language models Oct 28: Word embeddings
10	Nov 2, 4	Bayesian inference, topic models	<ul style="list-style-type: none">○ Mixtures○ Bayesian inference○ Topic models	Tue: Assn 4 in <ul style="list-style-type: none">○ Assn5 out	Nov 2: Bayesian inference Notes on Bayesian inference Nov 4: Bayes and topic models
11	Nov 9, 11	Introduction to neural networks	<ul style="list-style-type: none">○ Minimal neural network○ Regression examples	Thu: Assn 5 in <ul style="list-style-type: none">○ Assn6 out	Nov 9: Topic models Nov 11: Neural networks
12	Nov 16, 18	Deep neural networks	<ul style="list-style-type: none">○ Tensorflow playground○ Autoencoder examples	Tue: Quiz 3	Nov 16: Neural networks (continued) Notes on backpropagation Nov 18: Autoencoders
13	Nov 19-28	No class, Thanksgiving break			
14	Nov 30, Dec 2	Reinforcement learning	<ul style="list-style-type: none">○ Q-learning	Tue: Assn 6 in <ul style="list-style-type: none">○ Assn7 out	Nov 30: Reinforcement learning Dec 2: Deep reinforcement learning
15	Dec 7, 9	Societal issues for machine learning		Tue: Quiz 4	Dec 7: Societal issues

Many other topics we haven't touched

- Graphical models
- Collaborative filtering
- Generalized adversarial networks
- Kernel methods
- New types of generative models
- Causal inference and machine learning
- ...

Final exam

- Final exam Tuesday, December 21, 2021 at 7pm
- https://registrar.yale.edu/general-information/final-exams
- Review sessions, time and place TBA
- Length: About 1.5X Midterm
- Practice exams posted
- Any topic could be on exam...except

Vote a topic off the exam!



Nominations?

Your input

- Please complete a course review!
- I value your comments and feedback

Thank you!